

# How To Setup Grafana On Kubernetes

by **Bibin Wilson** · April 23, 2021



Grafana is an open-source lightweight dashboard tool. It can be integrated with many data sources like Prometheus, [AWS cloud watch](#), Stackdriver, etc. Running Grafana on Kubernetes

When Grafana is used with Prometheus, it caused PromQL to query metrics from Prometheus

In our previous posts, we have looked at the following.

- 2 [Setup Kube State Metrics](#)
- 3 [Setup alert manager on Kubernetes](#)

This tutorial explains how to run Grafana on [Kubernetes cluster](#). Using Grafana you can simplify Kubernetes monitoring dashboards from Prometheus metrics.

## Grafana Kubernetes Manifests

All the Kubernetes manifests (YAML files) used in this tutorial are [hosted on Github](#) as well. You can clone it and use it for the setup.

```
git clone https://github.com/bibinwilson/kubernetes-grafana.git
```

## Deploy Grafana On Kubernetes

Let’s look at the Grafana setup in detail.

**Step 1:** Create file named `grafana-datasource-config.yaml`

```
vi grafana-datasource-config.yaml
```

Copy the following contents.

**Note:** The following data source configuration is for Prometheus. If you have more data sources, you can add more data sources with different YAMLS under the data section.

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: grafana-datasources
  namespace: monitoring
data:
  prometheus.yaml: |-
    {
      "apiVersion": 1,
      "datasources": [
        {
          "access": "proxy",
          "editable": true,
          "name": "prometheus",
          "orgId": 1,
```

```
}  
]  
}
```

**Step 2:** Create the configmap using the following command.

```
kubectl create -f grafana-datasource-config.yaml
```

**Step 3:** Create a file named `deployment.yaml`

```
vi deployment.yaml
```

Copy the following contents on the file.

```
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: grafana  
  namespace: monitoring  
spec:  
  replicas: 1  
  selector:  
    matchLabels:  
      app: grafana  
  template:  
    metadata:  
      name: grafana  
      labels:  
        app: grafana  
    spec:  
      containers:  
        - name: grafana  
          image: grafana/grafana:latest  
          ports:  
            - name: grafana  
              containerPort: 3000  
          resources:  
            limits:  
              memory: "1Gi"  
              cpu: "1000m"  
            requests:  
              memory: 500M  
              cpu: "500m"  
          volumeMounts:  
            - mountPath: /var/lib/grafana  
              name: grafana-storage  
            - mountPath: /etc/grafana/provisioning/datasources  
              name: grafana-datasources  
              readOnly: false  
      volumes:  
        - name: grafana-storage  
          emptyDir: {}  
        - name: grafana-datasources  
          configMap:  
            defaultMode: 420  
            name: grafana-datasources
```

if you are deploying Grafana for your project requirements. It will persist all the configs and data that Grafana uses.

#### Step 4: Create the deployment

```
kubectl create -f deployment.yaml
```

#### Step 5: Create a service file named `service.yaml`

```
vi service.yaml
```

Copy the following contents. This will expose Grafana on `NodePort` 32000. You can also expose it [using ingress](#) or a Loadbalancer based on your requirement.

```
apiVersion: v1
kind: Service
metadata:
  name: grafana
  namespace: monitoring
  annotations:
    prometheus.io/scrape: 'true'
    prometheus.io/port: '3000'
spec:
  selector:
    app: grafana
  type: NodePort
  ports:
    - port: 3000
      targetPort: 3000
      nodePort: 32000
```

#### Step 6: Create the service.

```
kubectl create -f service.yaml
```

Now you should be able to access the Grafana dashboard **using any node IP** on port `32000` . Make sure the port is allowed in the firewall to be accessed from your workstation.

```
http://<your-node-ip>:32000
```

You can also use port forwarding using the following command.

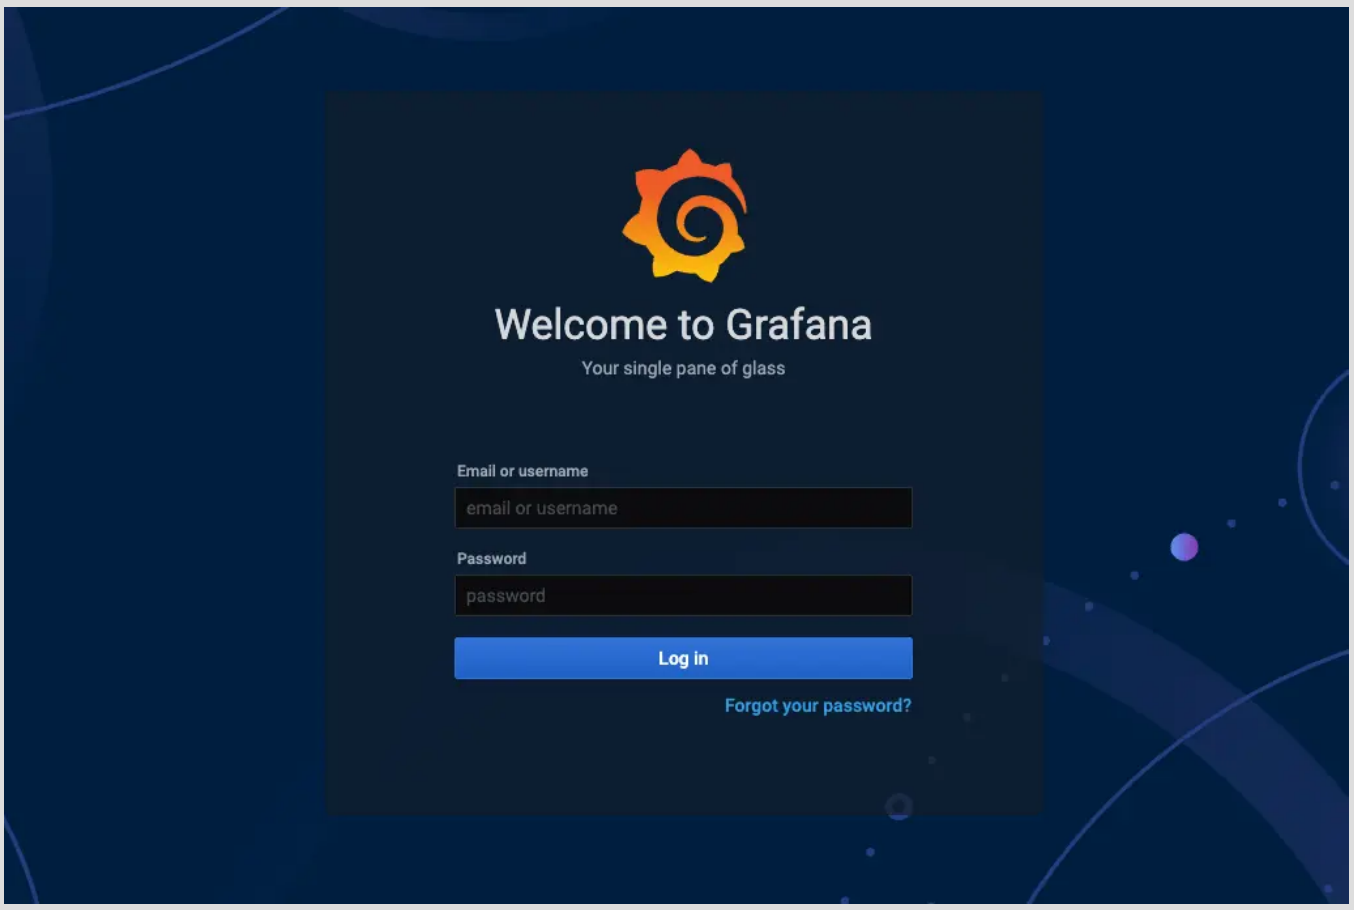
```
kubectl port-forward -n monitoring <grafana-pod-name> 3000 &
```

```
vagrant@dcubelab:~$ kubectl get po -n monitoring
NAME                                READY   STATUS    RESTARTS   AGE
grafana-64c89f57f7-kjqrb           1/1     Running   0           10m
vagrant@dcubelab:~$ kubectl port-forward -n monitoring grafana-64c89f57f7-kjqrb
3000 &
```

You will be able to access Grafana a from `http://localhost:3000`

Use the following default username and password to log in. Once you log in with default credentials, it will prompt you to change the default password.

```
User: admin
Pass: admin
```

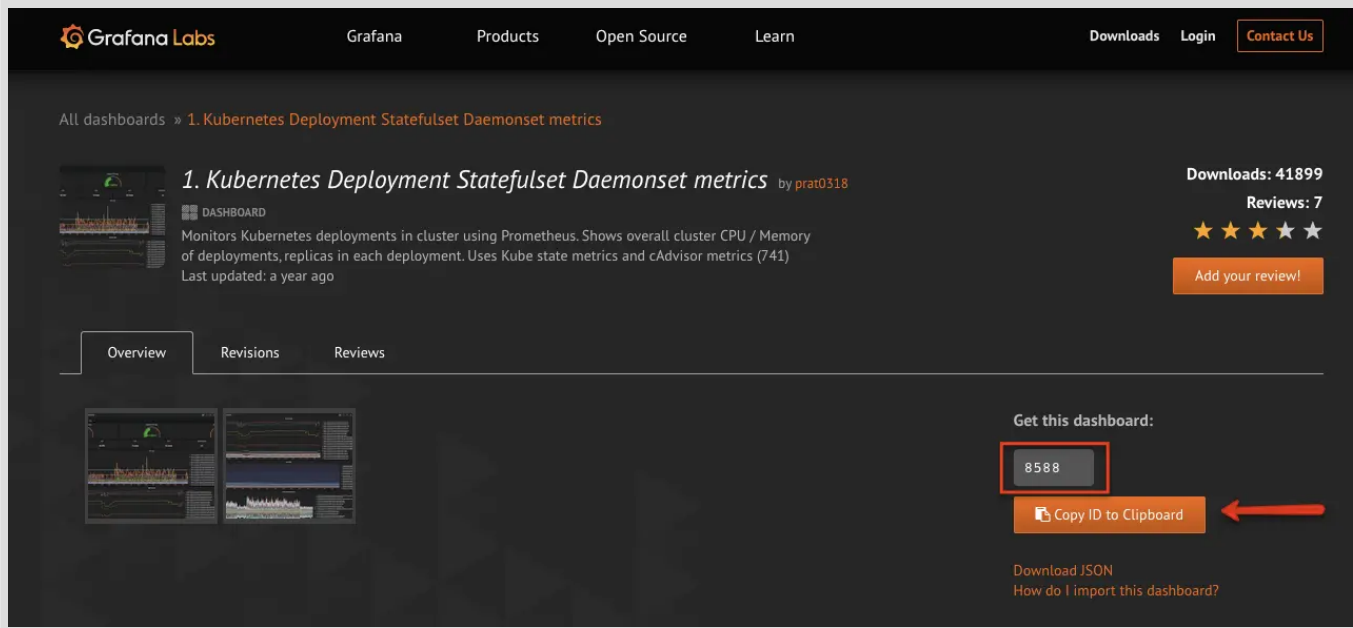


# Setup Kubernetes Dashboards on Grafana

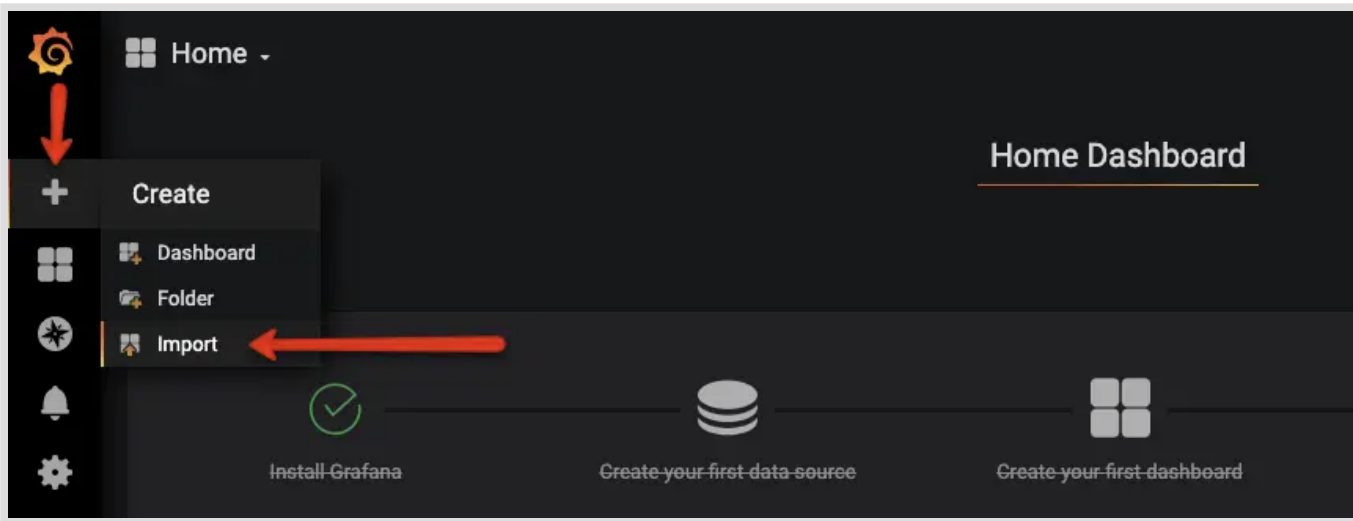
There are many prebuilt Grafana templates available for Kubernetes. To know more, see [Grafana Kubernetes Dashboard templates](#)

Setting up a dashboard from a template is pretty easy. Follow the steps given below to set up a Grafana dashboard to monitor kubernetes deployments.

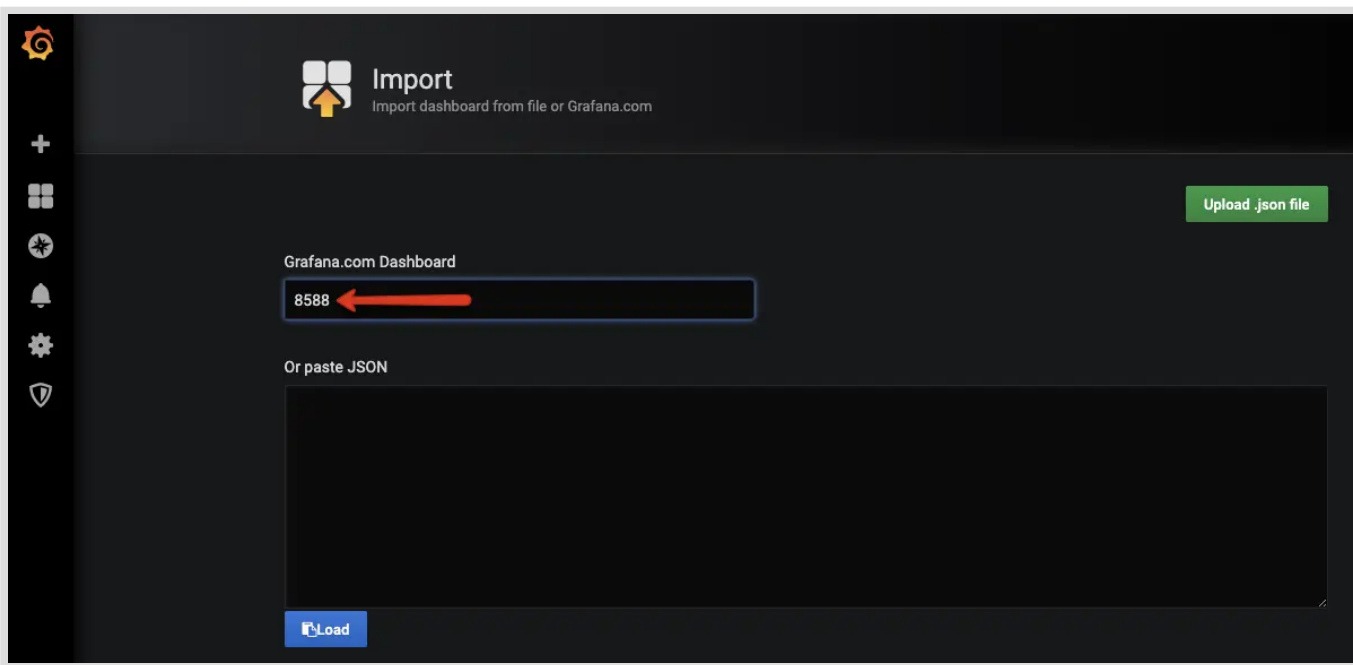
**Step 1:** Get the template ID from [grafana public template](#). as shown below.



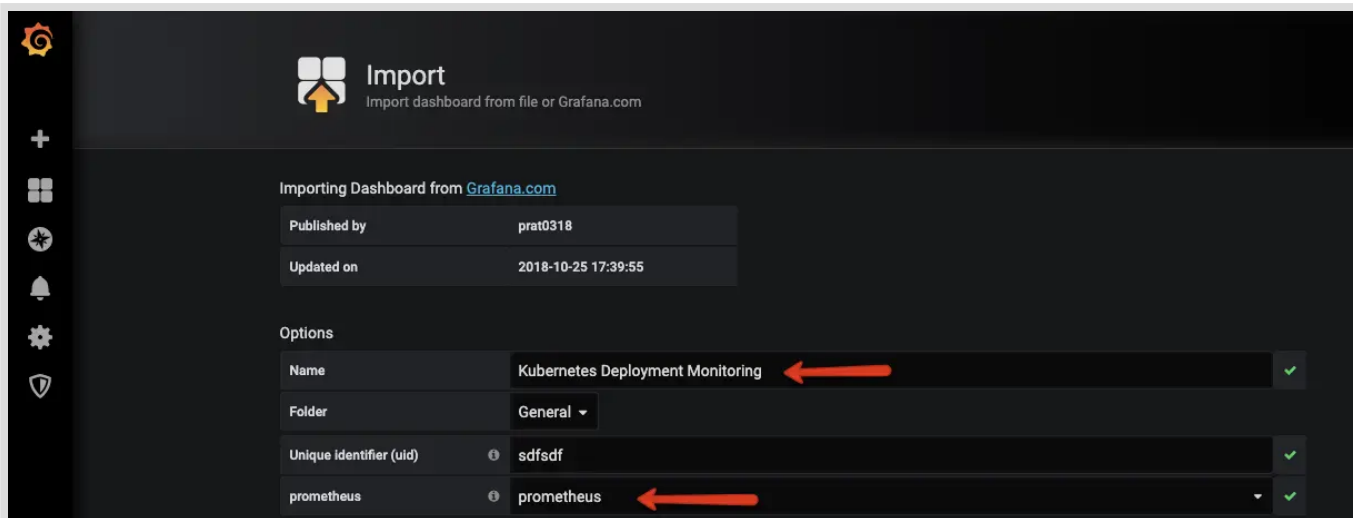
**Step 2:** Head over to grafana and select the import option.



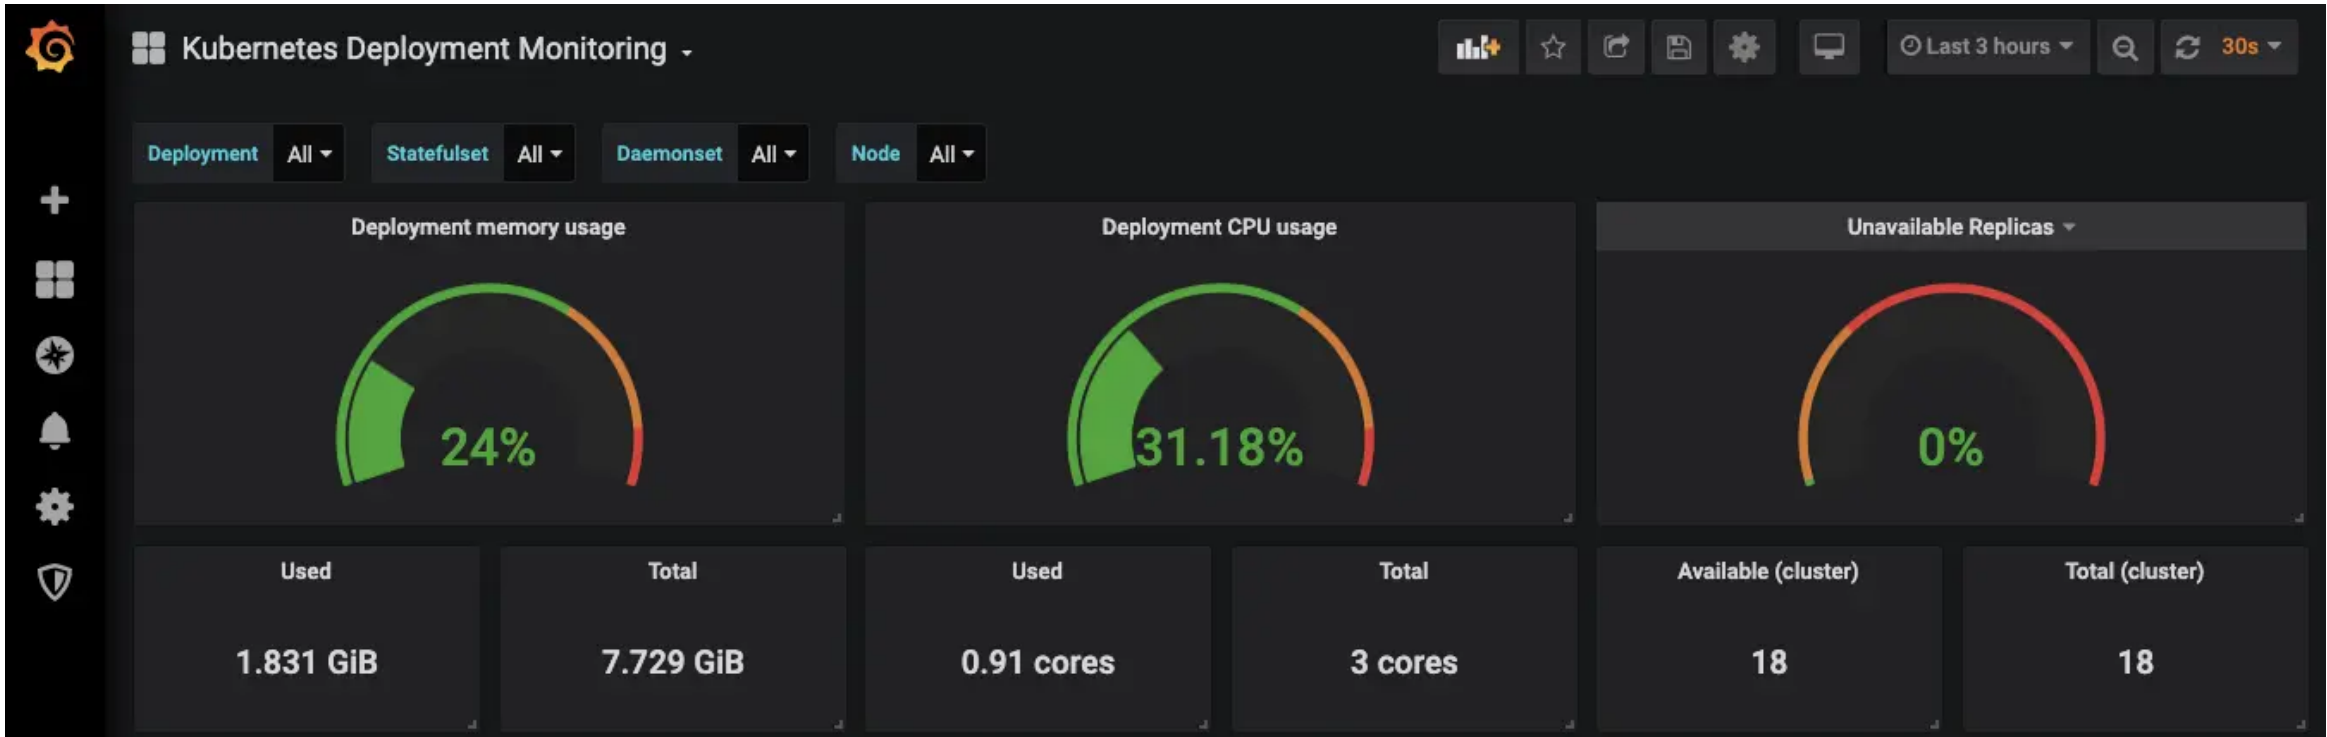
**Step 3:** Enter the dashboard ID you got it step 1



**Step 4:** Grafana will automatically fetch the template from Grafana website. You can change the values as shown in the image below and click import.



You should see the dashboard immediately.



## Conclusion

Grafana is a very powerful tool when it comes to monitoring dashboards.

It is used by many organizations to monitor their workloads.

Let us know how you are using Grafana in your organization. Also, let us know if you want to add more information to this article.

4  
SHARES



**TAGS:** Grafana Kubernetes Prometheus

**4 SHARES:** [f SHARE 4](#) [TWEET](#) [✉](#) [in](#)




**Bibin Wilson**

An author, blogger and DevOps practitioner. In spare time, he loves to try out the latest open source technologies. He works as an Associate Technical Architect



[VIEW COMMENTS \(32\)](#)

YOU MAY ALSO LIKE



**D** — DEVOPS



[Learn](#) [Resources](#) [News](#) [Newsletter](#) [Certification Guides](#)

[START HERE](#)







by **devopscube** · April 6, 2021

If you want to know how the Kubernetes nodes perform or monitor system-level insights of kubernetes nodes, you...



**D** — DEVOPS

## How To Create Jenkins Shared Library For Pipelines

by **devopscube** · July 7, 2019

In this tutorial, you will learn how to create a basic shared library and integrate it with Jenkins...



**K** — KUBERNETES

## Kubernetes monitoring With Sensu: Setting up Container Sidecar Agent

by **Jef Spaleta** · August 27, 2019

The rise of containerized infrastructure has caused us to rethink the way we build and deploy our applications....



**K** — KUBERNETES

## How To Setup Kubernetes Cluster On Google Cloud (GKE)

by **Bibin Wilson** · June 13, 2021

This guide walks you through deploying a Kubernetes Cluster on google cloud using the Google Kubernetes Engine (GKE)....



**D** — DEVOPS



[Learn](#) ▾ [Resources](#) ▾ [News](#) [Newsletter](#) [Certification Guides](#) ▾

START HERE ↗







# – Managed Cloud FileStore Service

by **devopscube** · August 16, 2018

Google Cloud Filestore is a managed NFS implementation on google cloud. This is one of the awaited features...



**C** — CLOUD

## How to Setup a Replicated GlusterFS Cluster on AWS EC2

by **devopscube** · October 6, 2016

GlusterFS is one of the best open source distributed file systems. If you want a highly available distributed...

DevopsCube

©devopscube 2021. All rights reserved.

[Privacy Policy](#)

[About](#)

[Site Map](#)

[Disclaimer](#)

[Contribute](#)

[Advertise](#)

[Archives](#)

