




How To Setup Jenkins On Kubernetes Cluster – Beginners Guide

by **Bibin Wilson** · May 9, 2021



Hosting Jenkins on a [Kubernetes cluster](#) is beneficial for Kubernetes-based deployments and dynamic container-based scalable Jenkins agents.

In this guide, I have explained the step-by-step process for setting up Jenkins on a [Kubernetes cluster](#).

Setup Jenkins On Kubernetes Cluster

For setting up a [Jenkins](#) cluster on Kubernetes, we will do the following.

3. Create local persistent volume for persistent Jenkins data on pod restarts.
4. Create a deployment YAML and deploy it.
5. Create a service YAML and deploy it.
6. Access the Jenkins application on a Node Port.

Note: This tutorial doesn't use local persistent volume as this is a generic guide. For using persistent volume for your Jenkins data, you need to create volumes of relevant cloud or on-prem data center and configure it.

Jenkins Kubernetes Manifest Files

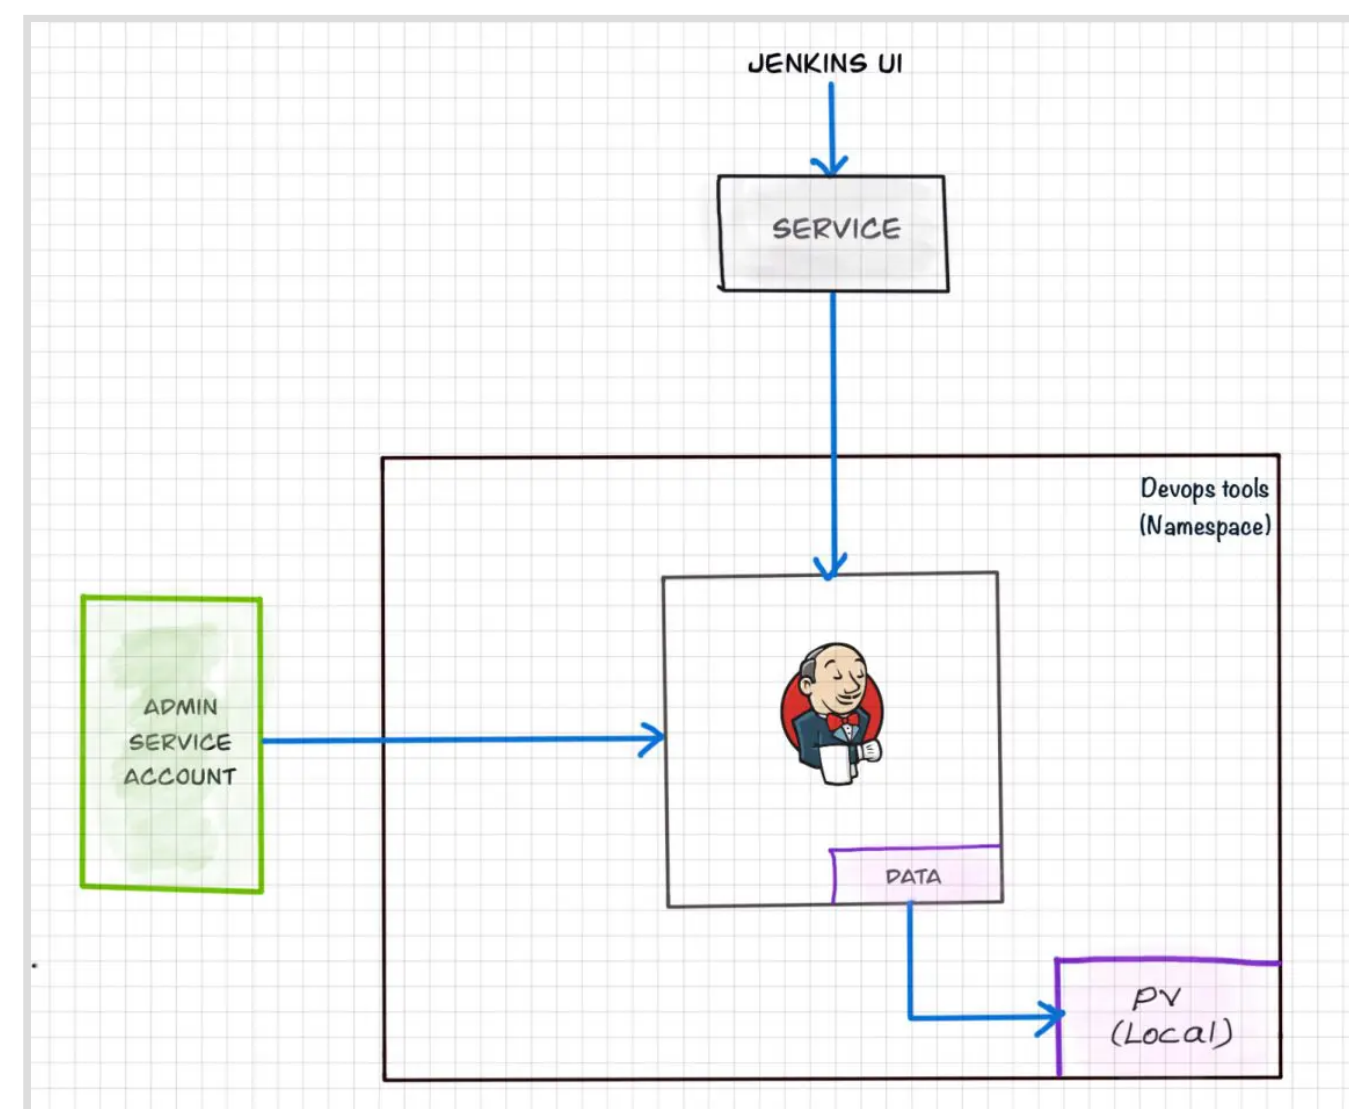
All the Jenkins Kubernetes manifest files used in this blog are [hosted on Github](#). Please clone the repository if you have trouble copying the manifest from the blog.

```
git clone https://github.com/scriptcamp/kubernetes-jenkins
```

Use the Github files for reference and follow the steps in the next sections

Kubernetes Jenkins Deployment

Here is the high level view of what we are going to do.



Step 1: Create a Namespace for Jenkins. It is good to categorize all the [devops tools](#) as a separate namespace from other applications.

```
kubectl create namespace devops-tools
```

Step 2: Create a `serviceAccount.yaml` file and copy the following admin service account manifest.

```
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: jenkins-admin
rules:
  - apiGroups: [""]
    resources: ["*"]
    verbs: ["*"]

---
apiVersion: v1
kind: ServiceAccount
metadata:
  name: jenkins-admin
  namespace: devops-tools

---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: jenkins-admin
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: jenkins-admin
subjects:
  - kind: ServiceAccount
    name: jenkins-admin
    namespace: devops-tools
```

The `serviceAccount.yaml` creates a `jenkins-admin` clusterRole, `jenkins-admin` [ServiceAccount](#) and binds the clusterRole to the service account.

The `jenkins-admin` cluster role has all the permissions to manage the cluster components. You can also restrict access by specifying individual resource actions.

Now create the service account using kubectl.

```
kubectl apply -f serviceAccount.yaml
```

Step 3: Create a `volume.yaml` and copy the following persistent volume manifest.

```
metadata:
  name: local-storage
provisioner: kubernetes.io/no-provisioner
volumeBindingMode: WaitForFirstConsumer

---
apiVersion: v1
kind: PersistentVolume
metadata:
  name: jenkins-pv-volume
  labels:
    type: local
spec:
  storageClassName: local-storage
  claimRef:
    name: jenkins-pv-claim
    namespace: devops-tools
  capacity:
    storage: 10Gi
  accessModes:
    - ReadWriteOnce
  local:
    path: /mnt
  nodeAffinity:
    required:
      nodeSelectorTerms:
        - matchExpressions:
            - key: kubernetes.io/hostname
              operator: In
              values:
                - worker-node01

---
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: jenkins-pv-claim
  namespace: devops-tools
spec:
  storageClassName: local-storage
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 3Gi
```

Important Note: Replace `worker-node01` with any one of your cluster worker nodes hostname.

You can get the worker node hostname using the kubectl.

```
kubectl get nodes
```

```
+ ~ kubectl get nodes
NAME          STATUS    ROLES          AGE   VERSION
master-node   Ready     control-plane,master  3d19h v1.21.0
worker-node01 Ready     <none>          3d19h v1.21.0
worker-node02 Ready     <none>          3d19h v1.21.0
```

node under `/mnt` location.

As the `local` storage class requires the node selector, you need to specify the worker node name correctly for the Jenkins pod to get scheduled in the specific node.

If the pod gets deleted or restarted, the data will get persisted in the node volume. However, if the node gets deleted, you will lose all the data.

Ideally, you should use a persistent volume using the available storage class with the cloud provider or the one provided by the cluster administrator to persist data on node failures.

Lets create the volume using kubectl

```
kubectl create -f volume.yaml
```

Step 2: Create a Deployment file named `deployment.yaml` and copy the following deployment manifest.

Here we are using the latest [Jenkins LTS docker image](#) from the Docker hub.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: jenkins
  namespace: devops-tools
spec:
  replicas: 1
  selector:
    matchLabels:
      app: jenkins-server
  template:
    metadata:
      labels:
        app: jenkins-server
    spec:
      securityContext:
        fsGroup: 1000
        runAsUser: 1000
      serviceAccountName: jenkins-admin
      containers:
        - name: jenkins
          image: jenkins/jenkins:lts
          resources:
            limits:
              memory: "2Gi"
              cpu: "1000m"
```

```

      cpu: 500m
    ports:
      - name: httpport
        containerPort: 8080
      - name: jnlpport
        containerPort: 50000
    livenessProbe:
      httpGet:
        path: "/login"
        port: 8080
      initialDelaySeconds: 90
      periodSeconds: 10
      timeoutSeconds: 5
      failureThreshold: 5
    readinessProbe:
      httpGet:
        path: "/login"
        port: 8080
      initialDelaySeconds: 60
      periodSeconds: 10
      timeoutSeconds: 5
      failureThreshold: 3
    volumeMounts:
      - name: jenkins-data
        mountPath: /var/jenkins_home
  volumes:
    - name: jenkins-data
      persistentVolumeClaim:
        claimName: jenkins-pv-claim

```

In this Jenkins [Kubernetes deployment](#) we have used the following.

- 1 `securityContext` for Jenkins pod to be able to write to the local persistent volume.
- 2 Liveliness and readiness probe.
- 3 Local persistent volume based on local storage class that holds the Jenkins data path `/var/jenkins_home`

Note: The deployment file uses local storage class persistent volume for Jenkins data. For production use cases, you should add a cloud-specific storage class persistent volume for your Jenkins data. See the sample implementation of [persistent volume for Jenkins in Google Kubernetes Engine](#)

If you don't want the local storage persistent volume, you can replace the volume definition in the deployment with the host directory as shown below.

```
emptyDir: {}
```

Create the deployment using kubectl.

```
kubectl apply -f deployment.yaml
```

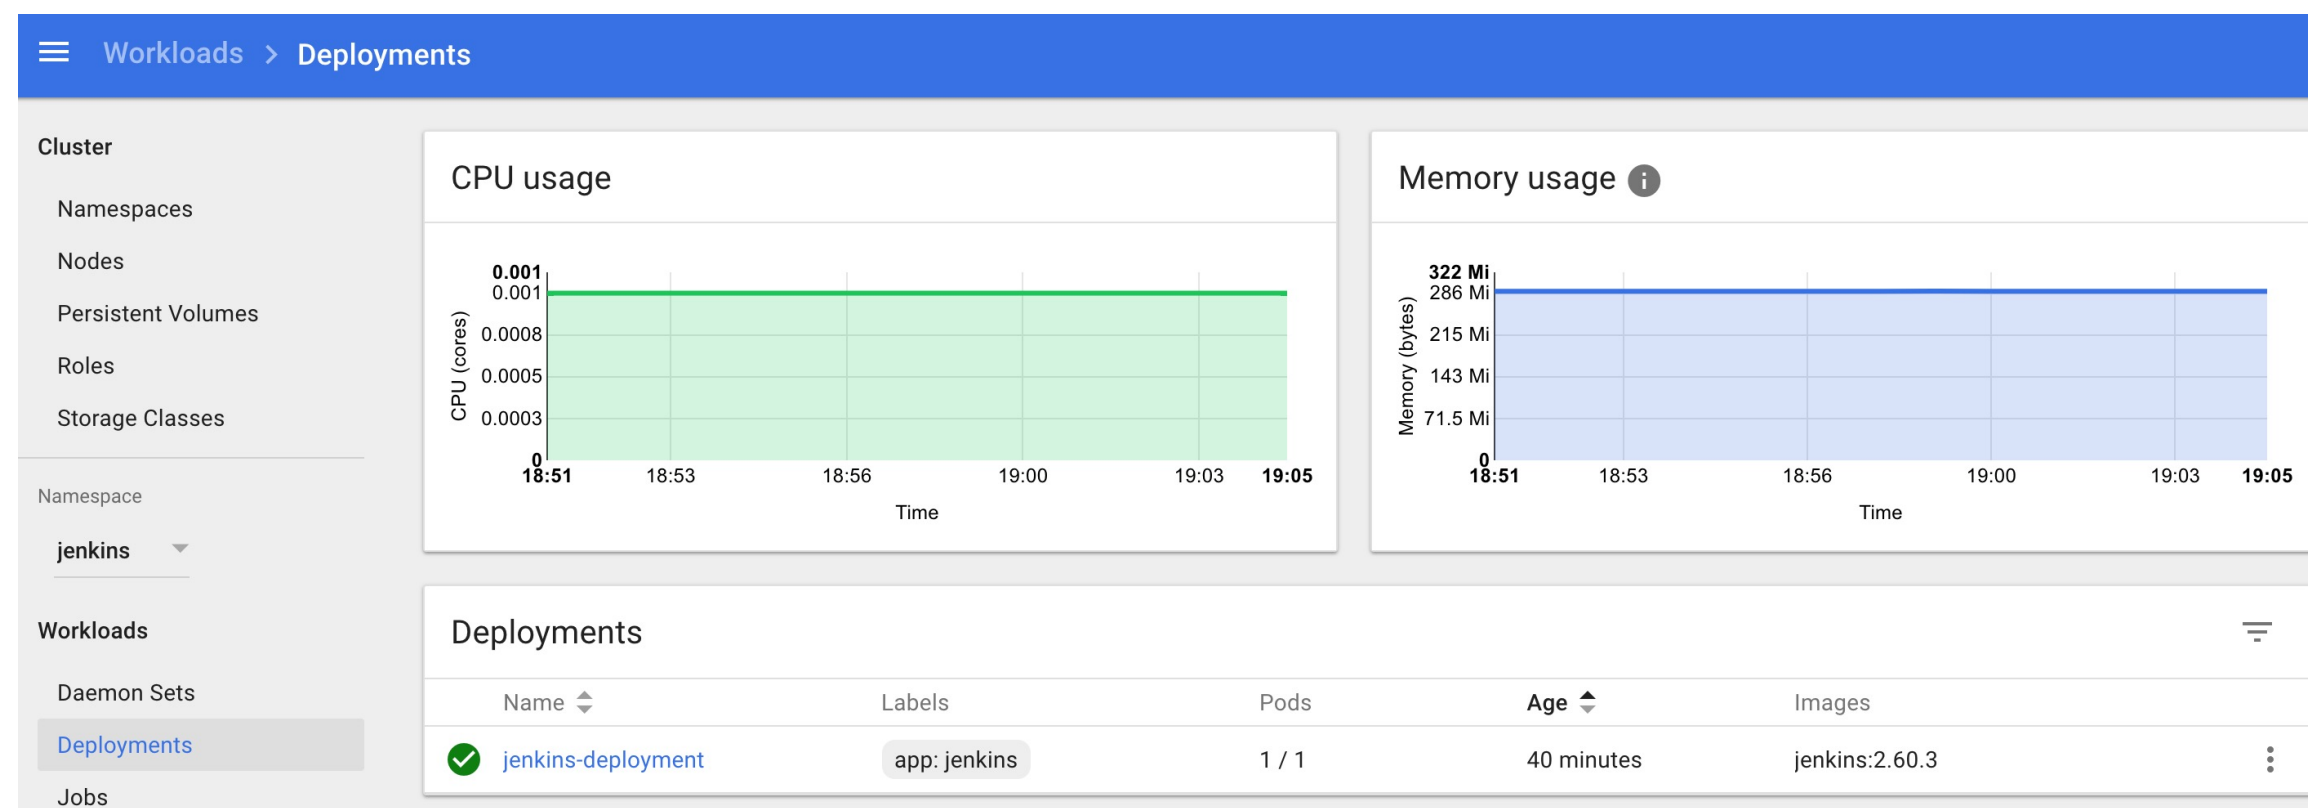
Check the deployment status.

```
kubectl get deployments -n devops-tools
```

Now, you can get the deployment details using the following command.

```
kubectl describe deployments --namespace=jenkins
```

Also, You can get the details from the kubernetes dashboard as shown below.



Accessing Jenkins Using Kubernetes Service

We have created a deployment. However, it is not accessible to the outside world. For accessing the Jenkins deployment from the outside world, we should create a service and map it to the deployment.

Step 1: Create a `service.yaml` and copy the following service manifest.

```
apiVersion: v1
kind: Service
metadata:
```

```

annotations:
  prometheus.io/scrape: 'true'
  prometheus.io/path:    /
  prometheus.io/port:    '8080'
spec:
  selector:
    app: jenkins-server
  type: NodePort
  ports:
    - port: 8080
      targetPort: 8080
      nodePort: 32000

```

Note: Here, we are using the type as `NodePort` which will expose Jenkins on all kubernetes node IPs on port 32000. If you have an [ingress setup](#), you can create an ingress rule to access Jenkins. Also, you can expose the Jenkins service as a Loadbalancer if you are running the cluster on AWS, Google, or Azure cloud.

Create the Jenkins service using kubectl.

```
kubectl apply -f service.yaml
```

Now if you browse to any one of the Node IPs on port `32000` , you will be able to access the Jenkins dashboard.

```
http://<node-ip>:32000
```

Jenkins will ask for the initial Admin password when you access the dashbaord for the first time.

Unlock Jenkins

To ensure Jenkins is securely set up by the administrator, a password has been written to the log (**not sure where to find it?**) and this file on the server:

```
/var/jenkins_home/secrets/initialAdminPassword
```

Please copy the password from either location and paste it below.

Administrator password

Continue

You can get that from the pod logs either from the kubernetes dashboard or CLI. You can get the pod details using the following CLI command.

```
kubectl get pods --namespace=devops-tools
```

And with the pod name, you can get the logs as shown below. replace the pod name with your pod name.

```
kubectl logs jenkins-deployment-2539456353-j00w5 --namespace=jenkins
```

The password can be found at the end of the log as shown below.

```
*****
*****
*****

Jenkins initial setup is required. An admin user has been created and a password generated.
Please use the following password to proceed to installation:

0d3a28914310427ea9c4c1e72554a238 ←
This may also be found at: /var/jenkins_home/secrets/initialAdminPassword
*****
*****
*****
```

Alternatively, you can run the exec command to get the password directly from the location as show below.

```
kubectl exec -it jenkins-559d8cd85c-cfcgk cat
/var/jenkins_home/secrets/initialAdminPassword -n devops-tools
```

dashboard.

Conclusion

When you host Jenkins on Kubernetes for production workloads, you need to consider setting up a highly available persistent volume to avoid data loss during pod or or node deletetion.

A pod or node deletion could happen anytime in Kubernetes environments. It could be a patching activity or a downscaling activity.

Hope this step by step guide helps you to learn and understand the components involved in setting up a Jenkins server on a Kubernetes cluster.

13 SHARES:

 SHARE 13

 TWEET








Bibin Wilson

An author, blogger and DevOps practitioner. In spare time, he loves to try out the latest open source technologies. He works as an Associate Technical Architect

[VIEW COMMENTS \(3\)](#)

YOU MAY ALSO LIKE



D — DEVOPS


How To Install Latest Sonatype Nexus 3 on Linux

by **devopscube** · April 25, 2021

Sonatype Nexus is one of the best open-source artifact management tools. It is some tool that you cannot...

K — KUBERNETES

How to Create kubernetes Role for



by **Bibin Wilson** · June 1, 2021

In this blog, you will learn how to create Kubernetes role for a service account and use it...




K — KUBERNETES

Kubernetes monitoring With Sensu: Setting up Container Sidecar Agent

by **Jef Spaleta** · August 27, 2019

The rise of containerized infrastructure has caused us to rethink the way we build and deploy our applications....




K — KUBERNETES

How To Setup Kubernetes Cluster Using Kubeadm

by **Bibin Wilson** · May 18, 2021

Kubeadm is an excellent tool to set up a working kubernetes cluster in less time. It does all...



K — KUBERNETES


How To Build Docker Image In Kubernetes Pod

by **devopscube** · July 24, 2021

This beginner's guide focuses on step by step process of setting up Docker image build in Kubernetes pod...



D — DEVOPS



GitHub

by **Julien Delange** · January 19, 2020

Code reviews are part of the daily activities of software engineers and a key process in release management....